

# Foundations of NumPy and Matplotlib

An Applied Python Mini-Chapter

Dr. Youssef Salman

**Learning outcomes.** By the end of this chapter, you will be able to:

- create and manipulate NumPy arrays (vectors and matrices);
- use vectorized operations and efficient slicing;
- generate synthetic data (linear grids, random samples, basic distributions);
- produce basic plots with Matplotlib.

## 0. Getting set up (Python environment)

**Prerequisites.** Basic Python (variables, lists, loops, functions) and an environment such as Jupyter (or Spyder/VS Code).

Quick install (terminal):

```
1 pip install --upgrade pip
2 pip install numpy matplotlib
```

Conventional imports in Python:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
```

## 1 Meet NumPy: the ndarray

**NumPy array (ndarray).** A contiguous, fixed-size, homogeneous data container (all elements share the same type). Enables fast vectorized operations and broadcasting, often much faster than pure-Python loops.

### 1.1 Creation and basic properties

```
1 import numpy as np
2
3 # 1) From a Python list
4 a = np.array([1, 1.5, 2, 2.5])
5 print(a.dtype, a.shape, a.ndim) # float64 (4,) 1
6
7 # 2) Common generators
```

```

8 z = np.zeros(5)           # [0., 0., 0., 0., 0.]
9 o = np.ones((2, 3))       # 2x3 matrix of ones
10 r = np.arange(0, 10, 2)    # 0, 2, 4, 6, 8
11 l = np.linspace(0, 1, 5)   # 5 points from 0 to 1 inclusive
12
13 # 3) Random numbers (seed for reproducibility)
14 rng = np.random.default_rng(seed=42)
15 u = rng.random(4)          # Uniform [0, 1)
16 n = rng.normal(0, 1, (2,2)) # Gaussian N(0,1) 2x2

```

**Key attributes.**

- `.shape`: tuple of dimensions, e.g. `(2, 3)` for a  $2 \times 3$  matrix.
- `.ndim`: number of dimensions (axes).
- `.dtype`: data type (e.g. `float64`, `int32`, `bool`).
- `.size`: total number of elements.

**Exercise (types and shapes).** Create an array of integers from  $-3$  to  $3$  included. Convert it to `float32`, then reshape it to a  $2 \times 3$  matrix (discard or adapt one element if needed). Verify `dtype`, `shape` and `size`.

## 1.2 Vectorized arithmetic & broadcasting

Vectorized operations apply elementwise without explicit loops, and *broadcasting* aligns shapes when possible:

```

1 x = np.array([0., 1., 2., 3.])
2 y = np.array([10., 10., 10., 10.])
3
4 s = x + y           # [10., 11., 12., 13.]
5 p = x * 2           # [0., 2., 4., 6.]
6 q = (x + 1) / (y)  # [0.1, 0.2, 0.3, 0.4]
7
8 # Broadcasting: adding a scalar or compatible shape
9 M = np.arange(6).reshape(2, 3)  # [[0,1,2],[3,4,5]]
10 v = np.array([1, 0, -1])
11 B = M + v          # [[1,1,1],[4,4,4]] (v broadcast across rows)

```

**Why it is fast.** NumPy offloads tight loops to optimized C/Fortran code and operates on contiguous memory blocks. Prefer vectorized code whenever possible.

## 1.3 Indexing and slicing

Slicing returns *views* whenever possible (no data copy), which is memory- and time-efficient.

```

1 a = np.arange(10)      # [0 1 2 3 4 5 6 7 8 9]
2 print(a[0], a[-1])    # first and last element
3 print(a[2:7])         # [2 3 4 5 6]

```

```

4 print(a[:5], a[5:])          # first five / from index 5 to end
5 print(a[::-2])              # step of 2: [0 2 4 6 8]
6
7 # 2D indexing (row, col)
8 M = np.arange(12).reshape(3,4)
9 print(M[0, 1])              # element at row 0, col 1
10 print(M[1, :])              # entire row 1 (view)
11 print(M[:, 2])              # entire column 2 (view)
12
13 # Modifying a view modifies the original!
14 row1 = M[1, :]
15 row1[:] = -1
16 print(M)                  # second row becomes [-1 -1 -1 -1]

```

**Exercise (views vs. copies).** Create a  $4 \times 4$  array with `np.arange`. Slice the central  $2 \times 2$  block into `C`, set all entries of `C` to 999, and inspect the original array. Then force a copy with `C = C.copy()` and repeat. What changes?

## 2 Array operations and reductions

### 2.1 Aggregations

NumPy provides fast aggregation functions that collapse an array to a scalar or along a given axis:

```

1 M = np.arange(1, 10).reshape(3, 3)
2
3 print(M.sum())              # total sum = 45
4 print(M.mean())             # average = 5.0
5 print(M.max(), M.min())    # 9, 1
6
7 # Axis-based reductions
8 print(M.sum(axis=0))       # column sums: [12 15 18]
9 print(M.sum(axis=1))       # row sums: [ 6 15 24]

```

**Exercise.** Create a random  $5 \times 5$  array from the standard normal distribution. Compute: mean of each row, standard deviation of each column. Then normalize the array by subtracting the global mean and dividing by the global standard deviation.

### 2.2 Boolean indexing and masks

Boolean arrays can be used to select or filter values:

```

1 x = np.arange(10)
2 mask = x % 2 == 0          # True for even numbers
3 print(x[mask])            # [0 2 4 6 8]

```

```

4
5 x[x > 5] = 99          # in-place modification
6 print(x)                # [ 0 1 2 3 4 5 99 99 99 99]

```

**Tip.** Boolean masks and vectorized comparisons are much faster and clearer than looping with `if` statements in Python.

## 3 Introduction to Matplotlib

### 3.1 Basic plotting

The most common interface is `matplotlib.pyplot`, which mimics MATLAB-style plotting.

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 x = np.linspace(0, 2*np.pi, 200)
5 y = np.sin(x)
6
7 plt.plot(x, y, label="sin(x)")
8 plt.xlabel("x")
9 plt.ylabel("y")
10 plt.title("Basic sine curve")
11 plt.legend()
12 plt.show()

```

**Exercise.** Plot both  $\sin(x)$  and  $\cos(x)$  on the same axes from 0 to  $2\pi$ . Use different line styles or colors, add a legend, and a grid.

### 3.2 Figures and subplots

Plots live inside a *figure*, which can contain multiple subplots arranged in a grid.

```

1 x = np.linspace(0, 1, 100)
2 y1 = x**2
3 y2 = np.sqrt(x)
4
5 fig, axes = plt.subplots(1, 2, figsize=(10, 4))
6
7 axes[0].plot(x, y1, color="red")
8 axes[0].set_title("y = x^2")
9
10 axes[1].plot(x, y2, color="blue")
11 axes[1].set_title("y = sqrt(x)")
12
13 plt.tight_layout()

```

```
14 plt.show()
```

**Note.** The object-oriented API (`fig, ax = plt.subplots(...)`) is recommended for complex plots, since it gives more explicit control.

### 3.3 Scatter plots and styles

```
1 rng = np.random.default_rng(0)
2 x = rng.normal(size=200)
3 y = rng.normal(size=200)
4
5 plt.scatter(x, y, alpha=0.6, c="purple", marker="x")
6 plt.title("Random scatter plot")
7 plt.xlabel("X")
8 plt.ylabel("Y")
9 plt.axis("equal")
10 plt.show()
```

**Exercise.** Generate 300 random points with coordinates  $(x, y)$  uniformly sampled in  $[-1, 1]^2$ . Plot them with a scatter plot, then highlight in a different color the points inside the unit circle  $x^2 + y^2 \leq 1$ .

## 4 Customizing plots

### 4.1 Colors, markers, and line styles

You can customize the appearance of curves with concise codes or full keywords.

```
1 x = np.linspace(0, 2*np.pi, 200)
2
3 plt.plot(x, np.sin(x), "r--", label="red dashed")
4 plt.plot(x, np.cos(x), color="green", marker="o",
5           linestyle=":", label="green dotted with circles")
6
7 plt.legend()
8 plt.title("Custom styles")
9 plt.show()
```

#### Style codes.

- Colors: 'r' red, 'b' blue, 'g' green, or named colors like "orange".
- Line styles: '-' solid, '--' dashed, ':' dotted, '-.' dash-dot.
- Markers: 'o' circle, 'x' cross, 's' square, etc.

## 4.2 Labels, grids, and legends

```

1 x = np.linspace(0, 10, 100)
2 y1 = np.exp(-0.1*x) * np.sin(x)
3 y2 = np.exp(-0.1*x) * np.cos(x)
4
5 plt.plot(x, y1, label="exp(-0.1x) * sin(x)")
6 plt.plot(x, y2, label="exp(-0.1x) * cos(x)")
7 plt.xlabel("x")
8 plt.ylabel("y")
9 plt.title("Damped oscillations")
10 plt.legend(loc="upper right")
11 plt.grid(True, linestyle="--", alpha=0.6)
12 plt.show()

```

## 4.3 Histograms

Histograms display the distribution of data.

```

1 rng = np.random.default_rng(123)
2 data = rng.normal(loc=0, scale=1, size=1000)
3
4 plt.hist(data, bins=30, color="skyblue",
5           edgecolor="black", alpha=0.7)
6 plt.title("Histogram of N(0,1)")
7 plt.xlabel("Value")
8 plt.ylabel("Frequency")
9 plt.show()

```

**Exercise.** Generate 10,000 samples from a uniform distribution on  $[0, 1]$ . Plot a histogram with 20 bins. Overlay a vertical line at the theoretical mean 0.5.

# 5 Practical session: Bringing NumPy and Matplotlib together

## 5.1 Example: projectile motion

We can simulate simple physics using NumPy arrays and visualize with Matplotlib.

```

1 g = 9.81      # gravity
2 v0 = 20.0     # initial speed (m/s)
3 theta = 45 * np.pi/180  # angle in radians
4
5 t = np.linspace(0, 3, 200)
6 x = v0 * np.cos(theta) * t
7 y = v0 * np.sin(theta) * t - 0.5*g*t**2
8
9 plt.plot(x, y)

```

```

10 plt.title("Projectile trajectory")
11 plt.xlabel("x (m)")
12 plt.ylabel("y (m)")
13 plt.axhline(0, color="black", linewidth=0.8)
14 plt.show()

```

**Exercise.** Modify the initial angle and velocity, and observe the changes in trajectory. Can you compute the maximum height and range analytically and compare with the plot?

## 5.2 Example: population growth (logistic model)

The logistic growth model is widely used in biology and computer science (e.g. spreading of information or viruses).

$$P(t) = \frac{K}{1 + Ae^{-rt}}$$

where  $K$  is the carrying capacity,  $r$  the growth rate, and  $A$  depends on the initial condition.

```

1 K = 1000      # carrying capacity
2 r = 0.5       # growth rate
3 P0 = 10        # initial population
4 A = (K - P0) / P0
5
6 t = np.linspace(0, 20, 200)
7 P = K / (1 + A * np.exp(-r*t))
8
9 plt.plot(t, P, label="Logistic growth")
10 plt.xlabel("Time")
11 plt.ylabel("Population")
12 plt.title("Logistic model")
13 plt.legend()
14 plt.show()

```

**Exercise.** Try different values of  $K$ ,  $r$ , and  $P_0$ . What happens if  $P_0$  is close to  $K$ ? What if  $r$  is very small?

## 6 Practical exercises

1. **Vector operations.** Create two random vectors of size 100. Compute:

- dot product,
- elementwise product,
- cosine similarity.

2. **Matrix manipulations.** Create a  $5 \times 5$  matrix with values  $1, 2, \dots, 25$ . Extract:

- the main diagonal,

- the last column,
- the submatrix of the four central elements.

3. **Visualization of functions.** Plot on the same figure:

$$f(x) = e^{-x^2}, \quad g(x) = \sin(5x) e^{-x^2}$$

on  $x \in [-3, 3]$ . Add labels, legend, and grid.

4. **Random data and histograms.** Generate 5000 samples from the exponential distribution with parameter  $\lambda = 2$ . Plot the histogram and compare the empirical mean with the theoretical mean  $1/\lambda$ .
5. **Real-world application (informatics).** Suppose a server records the number of requests per minute, modeled as a Poisson( $\lambda = 5$ ). Simulate 1000 minutes, plot the histogram, and discuss what the distribution suggests about server load.

**Tip.** Always label axes and use legends in plots. Visual clarity is as important as numerical accuracy.

## 7 Selected solutions and hints

### 1. Vector operations

```

1 rng = np.random.default_rng(0)
2 u = rng.random(100)
3 v = rng.random(100)
4
5 dot = np.dot(u, v)
6 elemwise = u * v
7 cos_sim = dot / (np.linalg.norm(u) * np.linalg.norm(v))
8
9 print("dot =", dot)
10 print("cosine similarity =", cos_sim)

```

### 2. Matrix manipulations

```

1 M = np.arange(1, 26).reshape(5, 5)
2 diag = np.diag(M)
3 last_col = M[:, -1]
4 central = M[1:4, 1:4]
5
6 print("Diagonal:", diag)
7 print("Last column:", last_col)
8 print("Central block:\n", central)

```

### 3. Visualization of functions

```

1 x = np.linspace(-3, 3, 400)
2 f = np.exp(-x**2)
3 g = np.sin(5*x) * np.exp(-x**2)
4
5 plt.plot(x, f, label="exp(-x^2)")
6 plt.plot(x, g, label="sin(5x) * exp(-x^2)")
7 plt.xlabel("x"); plt.ylabel("y")
8 plt.legend(); plt.grid(True)
9 plt.title("Function comparison")
10 plt.show()

```

## 4. Random data and histograms

```

1 lam = 2
2 samples = rng.exponential(1/lam, 5000)
3
4 plt.hist(samples, bins=40, color="lightcoral",
5           edgecolor="black", alpha=0.7, density=True)
6 plt.axvline(samples.mean(), color="blue", linestyle="--",
7              label=f"empirical mean {samples.mean():.2f}")
8 plt.axvline(1/lam, color="green", linestyle="-.",
9              label="theoretical mean 0.5")
10 plt.legend()
11 plt.title("Exponential distribution")
12 plt.show()

```

## 5. Real-world application: server load

```

1 lam = 5
2 data = rng.poisson(lam, 1000)
3
4 plt.hist(data, bins=range(0, 15), align="left",
5           rwidth=0.8, color="skyblue", edgecolor="black")
6 plt.xlabel("Requests per minute")
7 plt.ylabel("Frequency")
8 plt.title("Poisson simulation (lambda=5)")
9 plt.show()
10
11 print("Mean load =", data.mean())

```

**Interpretation.** The histogram is centered around 5, with most values between 2 and 8. This indicates moderate variability but stable average server load, consistent with a Poisson process.

## Conclusion

In this mini-chapter you have learned:

- how to create and manipulate arrays with NumPy;
- how vectorization, slicing, and broadcasting simplify code and improve performance;
- how to use Matplotlib for line plots, scatter plots, histograms, and subplots;
- how to connect mathematical models to data visualization in applied contexts.

**Next step.** In the following chapter we extend these foundations to linear algebra (matrix operations, decompositions) and their applications in computer science.