

# Linear Algebra for Machine Learning

**Youssef SALMAN**

Applied Mathematics  
Lebanese University — Faculty of Sciences (Section V)

**M1: Applied Mathematics & Statistics for Computational Sciences**

October 15, 2025

# Learning Objectives

- Understand linear maps, kernel/image, rank–nullity, and invertibility.
- Compute eigenvalues/eigenvectors; know when diagonalization is possible.
- Use SVD for low-rank approximation and PCA for dimensionality reduction.
- Implement these concepts in Python (NumPy/Scikit-Learn).

# Outline I

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## Definition 1

A **vector** in  $\mathbb{R}^d$  is an ordered list of  $d$  numbers:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}.$$

It represents a point or data item with  $d$  features.

### Norm (length):

$$\|\mathbf{x}\| = \sqrt{x_1^2 + x_2^2 + \cdots + x_d^2}.$$

### Dot product (similarity):

$$\mathbf{x} \cdot \mathbf{y} = \sum_i x_i y_i.$$

### Cosine similarity (angle-based):

$$\cos \theta = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}.$$

**Real-life:** In text mining, documents are vectors of word counts; cosine similarity finds which texts are most alike (search engines, recommender systems).

### Definition 2

Given vectors  $v_1, v_2, \dots, v_k \in \mathbb{R}^n$ , the **span** is

$$\text{span}\{v_1, v_2, \dots, v_k\} = \{\alpha_1 v_1 + \alpha_2 v_2 + \cdots + \alpha_k v_k \mid \alpha_i \in \mathbb{R}\}.$$

It is the set of all linear combinations of these vectors.

### Examples:

- In  $\mathbb{R}^2$ ,  $\text{span}\{(1, 0)\}$  is the  $x$ -axis.

- In  $\mathbb{R}^2$ ,  $\text{span}\{(1, 0), (0, 1)\} = \mathbb{R}^2$ .
- In  $\mathbb{R}^3$ ,  $\text{span}\{(1, 0, 0), (0, 1, 0)\}$  is the  $xy$ -plane.

**Geometric meaning:** The span describes the smallest subspace containing the given vectors (line, plane, or the whole space).

## Definition 3

A **matrix** is a rectangular array of numbers with  $m$  rows and  $n$  columns:

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}.$$

## Interpretations:

- Dataset: rows = observations, columns = features.
- Linear map:  $Ax$  transforms  $x \in \mathbb{R}^n$  into  $\mathbb{R}^m$ .

**Example:** In a dataset of 100 students and 5 grades, the data form a  $100 \times 5$  matrix. Multiplying by a weight vector computes each student's final score.

# Rank of a Matrix

- $\text{rank}(A) = \text{number of linearly independent columns (column rank).}$
- $\text{rank}(A) = \text{number of linearly independent rows (row rank).}$
- $\text{rank}(A) = \text{number of pivots in the row echelon form.}$

## Properties:

- $\text{rank}(A) \leq \min(m, n).$
- $\text{rank}(AB) \leq \min(\text{rank}(A), \text{rank}(B)).$
- $A$  is invertible  $\iff \text{rank}(A) = n$  (when  $A$  is square  $n \times n$ ).

**Geometric meaning:** Rank = the number of independent directions preserved by the transformation  $A$ .

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

# Linear Maps I

## Definition 4

A mapping  $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$  is called *linear* if it preserves **addition** and **scaling**:

$$T(\alpha\mathbf{x} + \beta\mathbf{y}) = \alpha T(\mathbf{x}) + \beta T(\mathbf{y}) \quad \text{for all } \mathbf{x}, \mathbf{y} \in \mathbb{R}^n, \alpha, \beta \in \mathbb{R}.$$

## Note 1

Every linear map can be represented by a matrix  $A \in \mathbb{R}^{m \times n}$ , so that

$$T(\mathbf{x}) = A\mathbf{x}.$$

This makes linear maps easy to compute on a computer.

**Example:** A  $90^\circ$  rotation in the plane is linear:

$$R(x, y) = (-y, x), \quad R(\mathbf{x}) = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix} \mathbf{x}.$$

**Real-life:** In image processing, rotating or scaling an image are linear maps on the pixel coordinate vectors.

## Definition 5

For a linear map  $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ :

$$\ker(T) = \{x \in \mathbb{R}^n \mid T(x) = \mathbf{0}\}, \quad \text{Im}(T) = \{T(x) \mid x \in \mathbb{R}^n\} \subseteq \mathbb{R}^m.$$

## Definition 6

A linear map  $T$  (or matrix  $A$ ) is **injective** if

$$T(x_1) = T(x_2) \Rightarrow x_1 = x_2,$$

that is, it never sends two different vectors to the same image.

Equivalently,

$$\ker(A) = \{0\}.$$

## Definition 7

Surjective (onto). A linear map  $T$  (or matrix  $A$ ) is **surjective** if

$$\text{Im}(A) = \mathbb{R}^m,$$

meaning every vector in the target space is the image of some  $x$ .

## Proposition 1

- $T$  is **one-to-one (injective)**  $\iff \ker(T) = \{\mathbf{0}\}$ .
- $T$  is **onto (surjective)**  $\iff \text{Im}(T) = \mathbb{R}^m$ .

**Example 1 (Kernel).** Let

$$A = \begin{bmatrix} 1 & 2 \\ 2 & 4 \end{bmatrix}, \quad T(x) = Ax.$$

Here  $\ker(T) = \text{span}\{(-2, 1)\}$ , a line in  $\mathbb{R}^2$ .  $\Rightarrow$  Many different inputs collapse to  $\mathbf{0}$ , hence not injective.

**Example 2 (Image).** Same  $A$ : every output lies on the line  $\text{span}\{(1, 2)\} \subset \mathbb{R}^2$ . So  $\text{Im}(T)$  is 1D, not the whole  $\mathbb{R}^2 \Rightarrow$  not onto.

**Real-life analogy:**

- $\ker(T)$  = information lost. (e.g. projecting a photo onto the  $x$ -axis deletes vertical detail).
- $\text{Im}(T)$  = possible outputs. (e.g. grayscale images are projections of RGB, image lives in a lower-dimensional space).

## Theorem 8 (Rank–Nullity)

For  $A \in \mathbb{R}^{m \times n}$ ,

$$\text{rank}(A) + \dim(\ker A) = n.$$

## Proposition 2

If  $A \in \mathbb{R}^{n \times n}$ , then  $A$  is invertible

$$\iff \text{rank}(A) = n \iff \ker(A) = \{0\} \iff \det(A) \neq 0.$$

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## Definition 9

Let  $A \in \mathbb{R}^{n \times n}$ . A scalar  $\lambda \in \mathbb{R}$  is called an **eigenvalue** of  $A$  if there exists a nonzero vector  $\mathbf{v}$  such that

$$A\mathbf{v} = \lambda\mathbf{v}.$$

The vector  $\mathbf{v}$  is then an **eigenvector** associated with  $\lambda$ .

## Note 2

*Eigenvalues satisfy the **characteristic equation**:*

$$\det(A - \lambda I) = 0.$$

**Intuition:** An eigenvector is a special direction that  $A$  does not change (it may stretch or shrink, but does not rotate away). The eigenvalue tells us the factor of stretching.

**Example 1 (Scaling).**

$$A = \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}$$

- $(1, 0)$  is an eigenvector with eigenvalue  $\lambda = 2$ . -  $(0, 1)$  is an eigenvector with eigenvalue  $\lambda = 3$ .  $\Rightarrow A$  scales  $x$ -axis by 2 and  $y$ -axis by 3.

## Example 2 (Rotation).

$$R = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}$$

This matrix rotates vectors by  $90^\circ$ .  
- No real vector keeps its direction.  
 $\Rightarrow R$  has no real eigenvalues, but has complex eigenvalues  $\lambda = \pm i$ .

## Example 3 (Shear).

$$S = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$$

-  $(1, 0)$  is an eigenvector with eigenvalue 1.  
- Geometrically, shear keeps the  $x$ -axis fixed but shifts other directions.

## Applications in ML/CS:

- **PCA:** eigenvectors of the covariance matrix give the directions of maximum variance.
- **Google PageRank:** the dominant eigenvector of a stochastic matrix gives webpage importance.
- **Clustering:** spectral clustering uses eigenvectors of graph Laplacians.
- **Stability analysis:** signs of eigenvalues determine whether a system converges or diverges.

# Diagonalization

## Proposition 3

A matrix  $A \in \mathbb{R}^{n \times n}$  is **diagonalizable** if there exists a basis of  $\mathbb{R}^n$  consisting of eigenvectors of  $A$ . Equivalently,  $A$  is diagonalizable  $\iff$  the sum of the dimensions of its eigenspaces equals  $n$ .

**Meaning:** Diagonalization means that  $A$  can be rewritten as

$$A = PDP^{-1},$$

where  $D$  is diagonal (containing eigenvalues) and the columns of  $P$  are eigenvectors. This representation simplifies computations such as powers of  $A$ , solving systems, or analyzing dynamics.

## Note 3

**Spectral Theorem:** Every real symmetric matrix is diagonalizable with an orthonormal eigenbasis. This ensures numerical stability and is the foundation for many algorithms in data science (e.g. PCA).



# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## SVD: Definition I

### Theorem 10 (Singular Value Decomposition)

For any  $A \in \mathbb{R}^{m \times n}$ , there exist

- an orthogonal matrix  $U \in \mathbb{R}^{m \times m}$ ,
- an orthogonal matrix  $V \in \mathbb{R}^{n \times n}$ ,
- a diagonal matrix  $\Sigma \in \mathbb{R}^{m \times n}$  with nonnegative entries,

such that

$$A = U\Sigma V^\top, \quad \Sigma = \text{diag}(\sigma_1, \dots, \sigma_r), \quad \sigma_1 \geq \dots \geq \sigma_r > 0.$$

**Interpretation:** The SVD expresses  $A$  as a product of three simple transformations:

- ①  $V^\top$ : rotates the input space,
- ②  $\Sigma$ : stretches or compresses along orthogonal axes (singular values),
- ③  $U$ : rotates into the output space.

### Note 4

The *singular values*  $\sigma_i$  are the square roots of the nonzero eigenvalues of  $A^T A$  (or  $AA^T$ ). They measure how much  $A$  stretches vectors in different directions.

# Best Low-Rank Approximation

Theorem 11 (Eckart–Young–Mirsky)

Given the SVD

$$A = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^\top \quad (r = \text{rank}(A)),$$

the best approximation of  $A$  by a matrix of rank  $k < r$  is

$$A_k = \sum_{i=1}^k \sigma_i \mathbf{u}_i \mathbf{v}_i^\top,$$

in both the spectral norm ( $\|\cdot\|_2$ ) and Frobenius norm.

**Meaning:**

- $A_k$  keeps only the  $k$  largest singular values.
- It is the “closest” rank- $k$  matrix to  $A$ .
- This underlies many applications (data compression, PCA, noise reduction).

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## Definition 12

**Principal Component Analysis (PCA)** finds new orthogonal directions (principal components) in the feature space that capture the maximum variance in the data. Projecting onto the first  $k$  components gives a  $k$ -dimensional representation that preserves as much variability (information) as possible.

**Geometric view:** PCA rotates the coordinate system so that:

- PC1 = direction of greatest variance.
- PC2 = next orthogonal direction of variance.
- Higher PCs capture progressively less variance.

**Machine learning motivation:**

- Reduce dimensionality while keeping essential information.
- Remove redundancy due to correlated features.
- Facilitate visualization of high-dimensional datasets.

## Note 5

If  $X \in \mathbb{R}^{n \times d}$  is a centered data matrix (rows = samples, columns = features):

- ① Compute covariance matrix  $C = \frac{1}{n-1} X^\top X$ .
- ② Solve eigenvalue problem  $Cv = \lambda v$ .
- ③ Sort eigenvectors by decreasing eigenvalues.
- ④ Keep the first  $k$  eigenvectors  $W_k = [v_1, \dots, v_k]$ .
- ⑤ Project data:  $Z = XW_k$ .

**Alternative:** Compute PCA directly by the SVD of  $X = U\Sigma V^\top$ . The right singular vectors (columns of  $V$ ) give the principal components, and singular values squared ( $\sigma_i^2$ ) give the variance explained.

## What PCA does:

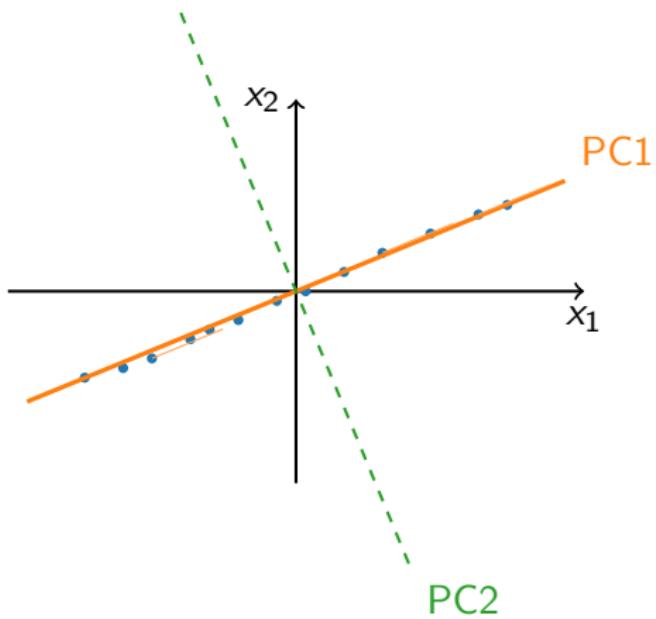
- Finds new coordinates aligned with variability in the data.
- Orders them by importance (variance explained).
- Allows truncation (keeping only first  $k$  PCs).

## Applications:

- Visualization of high-dimensional data (reduce to 2D or 3D).
- Noise reduction (discard low-variance components).
- Preprocessing for regression or clustering (avoid multicollinearity).
- Computer vision: eigenfaces for face recognition.
- Recommender systems: latent factors from user-item matrices.

**Key idea:** PCA = data compression without (much) loss of information.

# TikZ Illustration: PCA Projection (2D)



## Proposition 4

Let  $\lambda_1 \geq \dots \geq \lambda_d \geq 0$  be the eigenvalues of the covariance matrix. The proportion of variance captured by the first  $k$  principal components is

$$\text{Explained Variance Ratio}(k) = \frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i}.$$

### Interpretation:

- Each eigenvalue  $\lambda_i$  measures how much variance is along direction  $v_i$ .
- The ratio tells us how much information is kept when reducing to  $k$  dimensions.
- Example: If the first two PCs explain 90% of the variance, then projecting data to 2D keeps almost all the information.

### Use in practice:

- Helps decide how many PCs to keep (e.g., choose  $k$  so that  $> 95\%$  of variance is explained).
- Often visualized by a *scree plot*: eigenvalues or cumulative explained variance vs. number of components.

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## Definition 13

The **condition number** (in 2-norm) of a matrix  $A$  is

$$\kappa_2(A) = \frac{\sigma_{\max}(A)}{\sigma_{\min}(A)},$$

the ratio of the largest to the smallest singular value.

### Interpretation:

- If  $\kappa(A)$  is small (close to 1),  $A$  is *well-conditioned*: small changes in the data produce small changes in the solution.
- If  $\kappa(A)$  is large,  $A$  is *ill-conditioned*: even tiny data errors or rounding in the computer can lead to large errors in the result.

### Why it matters in PCA and ML:

- Covariance matrices with very different scales across features can have huge  $\kappa$ , making eigenvalue/SVD computations unstable.
- Multicollinearity (highly correlated features)  $\Rightarrow$  small singular values  $\Rightarrow$  large  $\kappa$ .
- This can cause unreliable PCs or regression coefficients.

## Note 6

### Practical tip:

- Always **center** features ( $\text{mean} = 0$ ).
- Often **scale** features ( $\text{variance} = 1$ ) to reduce ill-conditioning.
- Use **SVD-based PCA**, which is numerically more stable than directly solving the eigenvalue problem.

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

# NumPy: Eigenvalues & Eigenvectors

```
import numpy as np

A = np.array([[2., 1., 0.],
              [1., 2., 1.],
              [0., 1., 2.]])  
  
# Symmetric => use eigh (more stable)
w, V = np.linalg.eigh(A) # w: eigenvalues (asc), V: eigenvectors
idx = np.argsort(w)[::-1] # sort descending if desired
w, V = w[idx], V[:, idx]  
  
print("Eigenvalues:", w)
print("Orthonormal eigenvectors (columns of V):\n", V)
```

# NumPy: SVD & Best Rank- $k$ Approximation

```
import numpy as np

X = np.random.randn(100, 40) # data matrix (centered features recommended)
U, s, Vt = np.linalg.svd(X, full_matrices=False) #  $X = U @ np.diag(s) @ Vt$ 

k = 5
Xk = (U[:, :k] * s[:k]) @ Vt[:k, :]
err_fro = np.linalg.norm(X - Xk, 'fro') / np.linalg.norm(X, 'fro')

print(f"Relative Frobenius error (k={k}): {err_fro:.3f}")
```

## Scikit-Learn: PCA (Variance Explained)

```
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

# X: shape (n_samples, n_features)
X = np.random.randn(500, 50)

# Center & scale before PCA (common in practice)
Xz = StandardScaler(with_mean=True, with_std=True).fit_transform(X)

pca = PCA(n_components=0.95, svd_solver='full') # keep 95% variance
X_pca = pca.fit_transform(Xz)

print("Selected components:", pca.n_components_)
print("Explained variance ratio (first 10):", pca.explained_variance_ratio_[:10])
print("Cumulative:", np.cumsum(pca.explained_variance_ratio_)[:10])
```

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## Example 14

Let  $A = \begin{bmatrix} 1 & 2 & 3 \\ 0 & 1 & 1 \\ 1 & 1 & 2 \end{bmatrix}$ .

- Compute  $\text{rank}(A)$  via row-reduction or SVD.
- Find a basis for  $\ker(A)$ .
- Verify Rank–Nullity.

## Note 7

In Python, use `np.linalg.matrix_rank(A)` and solve  $Ax = \mathbf{0}$  with `scipy.linalg.null_space`.

## Example 15

Take the digits dataset (8x8 images) from `sklearn`. Apply PCA and visualize the first two PCs and reconstructions with  $k = 10, 20, 40$ .

## Note 8

*Use `sklearn.decomposition.PCA` and plot cumulative explained variance to select  $k$ .*

# Outline

- 1 Vectors and Matrices
- 2 Linear Maps, Kernel & Image
- 3 Eigenvalues, Diagonalization
- 4 Singular Value Decomposition (SVD)
- 5 Principal Component Analysis (PCA)
- 6 Numerical Stability & Conditioning
- 7 Python Implementations
- 8 Worked Examples
- 9 Conclusion & Takeaways

## Conclusion: Mathematical Takeaways

- **Vectors and Matrices:** building blocks for representing data and transformations.
- **Linear Maps:** connect matrices with geometric effects (scaling, rotation, projection).
- **Kernel & Image:** describe information lost and preserved by a transformation.
- **Rank:** measures the number of independent directions; links to invertibility.
- **Eigenvalues/Eigenvectors:** reveal invariant directions and scaling factors.
- **Matrix Decompositions (SVD):** universal tool for stability, compression, and analysis.

**Key idea:** Linear algebra provides the language and tools to manipulate high-dimensional data efficiently.

## Conclusion: Relevance to Machine Learning

- **Regression:** solution formulas involve inverses and matrix factorizations.
- **Dimensionality reduction:** PCA projects data onto principal components to reduce noise and redundancy.
- **Neural networks:** each layer applies linear transformations before nonlinear activations.
- **Data compression:** SVD and PCA approximate large datasets by low-rank structures.
- **Numerical stability:** conditioning and normalization are crucial for reliable algorithms.

**Take-home message:** Without linear algebra, machine learning cannot scale or even function. *This chapter builds the bridge from abstract math to practical AI.*

## Key Takeaways

- $\text{rank} + \text{nullity} = n$ ; invertibility  $\Leftrightarrow$  full rank.
- Eigen-decomposition diagonalizes symmetric matrices; use `eigh`.
- SVD exists for all matrices; gives optimal low-rank approximations.
- PCA = SVD on centered data; keep components by variance explained.
- Always consider conditioning and scaling for stable computations.

