

# Gradient Descent Methods: Full Batch, Stochastic and Mini-Batch

Youssef SALMAN

November 10, 2025

## Introduction

In optimization for Machine Learning, we often want to minimize a loss function

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n L(\theta; x_i),$$

where  $x_1, \dots, x_n$  are the training samples.

To minimize this loss, we use **gradient descent**. The update rule is always:

$$\theta_{k+1} = \theta_k - \eta g_k,$$

where:

- $\theta_k$  is the current parameter vector,
- $g_k$  is an estimate of the gradient at iteration  $k$ ,
- $\eta$  is the **learning rate**.

Depending on how we compute  $g_k$ , we obtain three versions of gradient descent.

## 1 1. Full Batch Gradient Descent

### Idea

Use **all** the dataset  $\{x_1, \dots, x_n\}$  to compute the gradient. This gives an exact and stable update.

### Mathematical Form

$$g_k = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} L(\theta_k; x_i).$$

## Advantages

- Very stable (no randomness).
- Exact direction of steepest descent.

## Disadvantages

- Very slow for large  $n$  (must compute gradient on all samples).
- Not suitable for big datasets.

## Steps (Algorithm)

1. Compute the full gradient:

$$g_k = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} L(\theta_k; x_i).$$

2. Update parameters:

$$\theta_{k+1} = \theta_k - \eta g_k.$$

3. Repeat until convergence.

## 2 2. Stochastic Gradient Descent (SGD)

### Idea

At each iteration, select **one random sample**  $z$  from the dataset, and compute the gradient only using this sample.

$$g_k = \nabla_{\theta} L(\theta_k; z).$$

This gives a very fast but noisy update.

### Why does this work?

Because:

$$\mathbb{E}[\nabla L(\theta; z)] = \frac{1}{n} \sum_{i=1}^n \nabla L(\theta; x_i)$$

so the average behavior matches the true gradient.

## Advantages

- Extremely fast updates.
- Works with huge datasets.
- Can escape shallow local minima.

## Disadvantages

- Very noisy.
- May oscillate around the minimum.

## Steps (Algorithm)

1. Randomly choose one sample:

$$z \sim \{x_1, \dots, x_n\}.$$

2. Compute stochastic gradient:

$$g_k = \nabla_{\theta} L(\theta_k; z).$$

3. Update:

$$\theta_{k+1} = \theta_k - \eta g_k.$$

4. Repeat.

**Example: what does “choose one sample” mean in SGD?** In Machine Learning, a sample is usually a **vector of features** (and possibly a label), not a single number. For example, consider the dataset:

$$x_1 = (2, 1, 3), \quad x_2 = (0, 4, 2), \quad x_3 = (1, 2, 5), \quad x_4 = (3, 0, 1).$$

At each iteration of SGD, we select one full observation at random:

$$z \sim \{x_1, x_2, x_3, x_4\}.$$

A possible sequence of chosen samples may be:

$$\begin{aligned} \text{Iteration 1: } z &= x_3 = (1, 2, 5), \\ \text{Iteration 2: } z &= x_1 = (2, 1, 3), \\ \text{Iteration 3: } z &= x_4 = (3, 0, 1), \\ \text{Iteration 4: } z &= x_2 = (0, 4, 2). \end{aligned}$$

The stochastic gradient at iteration  $k$  is computed using only this sample  $z$ :

$$g_k = \nabla_{\theta} L(\theta_k; z),$$

and the parameters are updated by

$$\theta_{k+1} = \theta_k - \eta g_k.$$

This process is repeated for many iterations (or epochs) until convergence.

**When do we stop?** Stochastic Gradient Descent does not stop after a single update. We repeat the procedure until a stopping criterion is satisfied. Common stopping conditions include:

- a fixed number of **epochs** (one epoch = a full pass through all  $n$  samples),
- the loss  $J(\theta)$  stops decreasing (convergence),
- the gradient norm becomes small:

$$\|\nabla J(\theta_k)\| < \varepsilon,$$

- a maximum number of iterations is reached.

In practice, SGD is usually run for several epochs because a single sample does not contain enough information to reach the minimum.

### 3 3. Mini-Batch Gradient Descent

#### Idea

At each iteration, randomly choose a **small subset** (a mini-batch) of size  $m$ :

$$B_k = \{x_{i_1}, \dots, x_{i_m}\}.$$

Then compute the average gradient over this mini-batch:

$$g_k = \frac{1}{m} \sum_{x \in B_k} \nabla_{\theta} L(\theta_k; x).$$

#### Why does mini-batch converge?

Even though each update uses only  $m < n$  samples, each sample has a chance to appear in many batches. On average:

$$\mathbb{E}[g_k] = \text{true gradient.}$$

Thus it converges (with proper learning rate).

#### Advantages

- Faster than full batch.
- More stable than SGD.
- Works very well with GPUs.

## Disadvantages

- Still contains some noise.
- Must choose batch size  $m$  (common choice: 32, 64, 128).

## Steps (Algorithm)

1. Randomly choose a mini-batch  $B_k$  of size  $m$ .

2. Compute mini-batch gradient:

$$g_k = \frac{1}{m} \sum_{x \in B_k} \nabla_{\theta} L(\theta_k; x).$$

3. Update:

$$\theta_{k+1} = \theta_k - \eta g_k.$$

4. Repeat.

**When do we stop?** Mini-Batch Gradient Descent also repeats the update many times. We continue until a stopping criterion is satisfied. Typical stopping conditions are:

- a fixed number of **epochs** (one epoch = using all  $n$  samples, divided into batches of size  $m$ ),
- the loss function  $J(\theta)$  stops decreasing,
- the gradient norm becomes small:

$$\|\nabla J(\theta_k)\| < \varepsilon,$$

- a maximum number of mini-batch iterations is reached.

Because each mini-batch only uses  $m < n$  samples, we must run several epochs. The random reshuffling of the data at each epoch ensures that every example contributes to the learning process.

## 4 4. The Learning Rate $\eta$

### Definition

The quantity  $\eta > 0$  is the **learning rate**. It controls how big the update step is:

$$\theta_{k+1} = \theta_k - \eta g_k.$$

## Role of $\eta$

- If  $\eta$  is too large: updates diverge or oscillate.
- If  $\eta$  is too small: convergence becomes extremely slow.
- Correct choice of  $\eta$  balances speed and stability.

## Interpretation

$\eta$  is the “speed” of learning.

A large  $\eta$  learns fast but may lose control. A small  $\eta$  learns safely but slowly.

## Typical Choices

- Batch GD: constant  $\eta$ .
- SGD: decreasing schedule  $\eta_k \rightarrow 0$ .
- Mini-batch: constant  $\eta$  (standard in deep learning).

## Summary Box

Batch uses all data. SGD uses 1 sample. Mini-batch uses  $m$  samples.

$\eta$  controls how big the step is: small  $\eta$  = slow, big  $\eta$  = unstable.